

Cyclistic Ride Share

Abdulhafiz Umar Dabo

2022-10-06

RIDE SHARE CASE STUDY: A PROJECT OF THE GOOGLE DATA ANALYTICS PROFESSIONAL CERTIFICATE

The goal of this project is to conduct analysis to help answer the key question: “In what ways do members and casual riders use Divvy bikes differently?” For comprehensive information about the project, Kindly check the README file in this repository.

SETTING UP THE ENVIRONMENT

We shall begin by loading the required packages onto the environment before importing the datasets.

```
library(tidyverse) #for data import and wrangling
```

```
## -- Attaching packages ----- tidyverse 1.3.2 --
## v ggplot2 3.3.6      v purrr   0.3.4
## v tibble  3.1.8      v dplyr   1.0.10
## v tidyr   1.2.1      v stringr 1.4.1
## v readr   2.1.2      v forcats 0.5.2
## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()    masks stats::lag()
```

```
library(lubridate) #helps wrangle date attributes
```

```
##
## Attaching package: 'lubridate'
##
## The following objects are masked from 'package:base':
##
##   date, intersect, setdiff, union
```

```
library(ggplot2) #for data visualization
```

IMPORT DATA

We then upload the Divvy datasets (csv files) here.

```
q2_2019 <- read_csv("Divvy_Trips_2019_Q2.csv")
```

```
## Rows: 1108163 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr  (4): 03 - Rental Start Station Name, 02 - Rental End Station Name, User...
## dbl  (5): 01 - Rental Details Rental ID, 01 - Rental Details Bike ID, 03 - R...
## dtm  (2): 01 - Rental Details Local Start Time, 01 - Rental Details Local En...
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
q3_2019 <- read_csv("Divvy_Trips_2019_Q3.csv")
```

```
## Rows: 1640718 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr  (4): from_station_name, to_station_name, usertype, gender
## dbl  (5): trip_id, bikeid, from_station_id, to_station_id, birthyear
## dtm  (2): start_time, end_time
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
q4_2019 <- read_csv("Divvy_Trips_2019_Q4.csv")
```

```
## Rows: 704054 Columns: 12
## -- Column specification -----
## Delimiter: ","
## chr  (4): from_station_name, to_station_name, usertype, gender
## dbl  (5): trip_id, bikeid, from_station_id, to_station_id, birthyear
## dtm  (2): start_time, end_time
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

```
q1_2020 <- read_csv("Divvy_Trips_2020_Q1.csv")
```

```
## Rows: 426887 Columns: 13
## -- Column specification -----
## Delimiter: ","
## chr  (5): ride_id, rideable_type, start_station_name, end_station_name, memb...
## dbl  (6): start_station_id, end_station_id, start_lat, start_lng, end_lat, e...
## dtm  (2): started_at, ended_at
##
## i Use 'spec()' to retrieve the full column specification for this data.
## i Specify the column types or set 'show_col_types = FALSE' to quiet this message.
```

WRANGLE DATA AND COMBINE INTO A SINGLE FILE

We first Compare column names in each of the files. While the names don't have to be in the same order, they DO need to match perfectly before we can use a command to join them into one file.

```
colnames(q3_2019)
```

```
## [1] "trip_id"          "start_time"       "end_time"
## [4] "bikeid"           "tripduration"     "from_station_id"
## [7] "from_station_name" "to_station_id"    "to_station_name"
## [10] "usertype"         "gender"           "birthyear"
```

```
colnames(q4_2019)
```

```
## [1] "trip_id"          "start_time"       "end_time"
## [4] "bikeid"           "tripduration"     "from_station_id"
## [7] "from_station_name" "to_station_id"    "to_station_name"
## [10] "usertype"         "gender"           "birthyear"
```

```
colnames(q2_2019)
```

```
## [1] "01 - Rental Details Rental ID"
## [2] "01 - Rental Details Local Start Time"
## [3] "01 - Rental Details Local End Time"
## [4] "01 - Rental Details Bike ID"
## [5] "01 - Rental Details Duration In Seconds Uncapped"
## [6] "03 - Rental Start Station ID"
## [7] "03 - Rental Start Station Name"
## [8] "02 - Rental End Station ID"
## [9] "02 - Rental End Station Name"
## [10] "User Type"
## [11] "Member Gender"
## [12] "05 - Member Details Member Birthday Year"
```

```
colnames(q1_2020)
```

```
## [1] "ride_id"          "rideable_type"    "started_at"
## [4] "ended_at"         "start_station_name" "start_station_id"
## [7] "end_station_name" "end_station_id"   "start_lat"
## [10] "start_lng"        "end_lat"          "end_lng"
## [13] "member_casual"
```

```
(q4_2019 <- rename(q4_2019
  ,ride_id = trip_id
  ,rideable_type = bikeid
  ,started_at = start_time
  ,ended_at = end_time
  ,start_station_name = from_station_name
  ,start_station_id = from_station_id
  ,end_station_name = to_station_name
  ,end_station_id = to_station_id
  ,member_casual = usertype))
```

Renaming columns to make them consistent with q1_2020 (as this will be the supposed going-forward table design for Divvy)

```
## # A tibble: 704,054 x 12
##   ride_id started_at      ended_at      rideable_t~1 tripd~2 start~3
##   <dbl> <dtm>          <dtm>          <dbl>    <dbl>    <dbl>
## 1 25223640 2019-10-01 00:01:39 2019-10-01 00:17:20      2215      940      20
## 2 25223641 2019-10-01 00:02:16 2019-10-01 00:06:34      6328      258      19
## 3 25223642 2019-10-01 00:04:32 2019-10-01 00:18:43      3003      850      84
## 4 25223643 2019-10-01 00:04:32 2019-10-01 00:43:43      3275     2350     313
## 5 25223644 2019-10-01 00:04:34 2019-10-01 00:35:42      5294     1867     210
## 6 25223645 2019-10-01 00:04:38 2019-10-01 00:10:51      1891      373     156
## 7 25223646 2019-10-01 00:04:52 2019-10-01 00:22:45      1061     1072      84
## 8 25223647 2019-10-01 00:04:57 2019-10-01 00:29:16      1274     1458     156
## 9 25223648 2019-10-01 00:05:20 2019-10-01 00:29:18      6011     1437     156
## 10 25223649 2019-10-01 00:05:20 2019-10-01 02:23:46      2957     8306     336
## # ... with 704,044 more rows, 6 more variables: start_station_name <chr>,
## #   end_station_id <dbl>, end_station_name <chr>, member_casual <chr>,
## #   gender <chr>, birthyear <dbl>, and abbreviated variable names
## #   1: rideable_type, 2: tripduration, 3: start_station_id
```

```
(q3_2019 <- rename(q3_2019
  ,ride_id = trip_id
  ,rideable_type = bikeid
  ,started_at = start_time
  ,ended_at = end_time
  ,start_station_name = from_station_name
  ,start_station_id = from_station_id
  ,end_station_name = to_station_name
  ,end_station_id = to_station_id
  ,member_casual = usertype))
```

```
## # A tibble: 1,640,718 x 12
##   ride_id started_at      ended_at      rideable_t~1 tripd~2 start~3
##   <dbl> <dtm>          <dtm>          <dbl>    <dbl>    <dbl>
## 1 23479388 2019-07-01 00:00:27 2019-07-01 00:20:41      3591     1214     117
## 2 23479389 2019-07-01 00:01:16 2019-07-01 00:18:44      5353     1048     381
## 3 23479390 2019-07-01 00:01:48 2019-07-01 00:27:42      6180     1554     313
## 4 23479391 2019-07-01 00:02:07 2019-07-01 00:27:10      5540     1503     313
## 5 23479392 2019-07-01 00:02:13 2019-07-01 00:22:26      6014     1213     168
## 6 23479393 2019-07-01 00:02:21 2019-07-01 00:07:31      4941      310     300
## 7 23479394 2019-07-01 00:02:24 2019-07-01 00:23:12      3770     1248     168
## 8 23479395 2019-07-01 00:02:26 2019-07-01 00:28:16      5442     1550     313
## 9 23479396 2019-07-01 00:02:34 2019-07-01 00:28:57      2957     1583      43
## 10 23479397 2019-07-01 00:02:45 2019-07-01 00:29:14      6091     1589      43
## # ... with 1,640,708 more rows, 6 more variables: start_station_name <chr>,
## #   end_station_id <dbl>, end_station_name <chr>, member_casual <chr>,
## #   gender <chr>, birthyear <dbl>, and abbreviated variable names
## #   1: rideable_type, 2: tripduration, 3: start_station_id
```

```
(q2_2019 <- rename(q2_2019
  ,ride_id = "01 - Rental Details Rental ID"
  ,rideable_type = "01 - Rental Details Bike ID"
```

```
,started_at = "01 - Rental Details Local Start Time"
,ended_at = "01 - Rental Details Local End Time"
,start_station_name = "03 - Rental Start Station Name"
,start_station_id = "03 - Rental Start Station ID"
,end_station_name = "02 - Rental End Station Name"
,end_station_id = "02 - Rental End Station ID"
,member_casual = "User Type"))
```

```
## # A tibble: 1,108,163 x 12
##   ride_id started_at ended_at rideable_t~1 01 - ~2 start~3
##   <dbl> <dtm> <dtm> <dbl> <dbl> <dbl>
## 1 22178529 2019-04-01 00:02:22 2019-04-01 00:09:48 6251 446 81
## 2 22178530 2019-04-01 00:03:02 2019-04-01 00:20:30 6226 1048 317
## 3 22178531 2019-04-01 00:11:07 2019-04-01 00:15:19 5649 252 283
## 4 22178532 2019-04-01 00:13:01 2019-04-01 00:18:58 4151 357 26
## 5 22178533 2019-04-01 00:19:26 2019-04-01 00:36:13 3270 1007 202
## 6 22178534 2019-04-01 00:19:39 2019-04-01 00:23:56 3123 257 420
## 7 22178535 2019-04-01 00:26:33 2019-04-01 00:35:41 6418 548 503
## 8 22178536 2019-04-01 00:29:48 2019-04-01 00:36:11 4513 383 260
## 9 22178537 2019-04-01 00:32:07 2019-04-01 01:07:44 3280 2137 211
## 10 22178538 2019-04-01 00:32:19 2019-04-01 01:07:39 5534 2120 211
## # ... with 1,108,153 more rows, 6 more variables: start_station_name <chr>,
## # end_station_id <dbl>, end_station_name <chr>, member_casual <chr>,
## # 'Member Gender' <chr>, '05 - Member Details Member Birthday Year' <dbl>,
## # and abbreviated variable names 1: rideable_type,
## # 2: '01 - Rental Details Duration In Seconds Uncapped', 3: start_station_id
```

```
str(q1_2020)
```

Inspecting the dataframes to look for incongruencies

```
## spec_tbl_df [426,887 x 13] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ ride_id : chr [1:426887] "EACB19130B0CDA4A" "8FED874C809DC021" "789F3C21E472CA96" "C9A3
## $ rideable_type : chr [1:426887] "docked_bike" "docked_bike" "docked_bike" "docked_bike" ...
## $ started_at : POSIXct[1:426887], format: "2020-01-21 20:06:59" "2020-01-30 14:22:39" ...
## $ ended_at : POSIXct[1:426887], format: "2020-01-21 20:14:30" "2020-01-30 14:26:22" ...
## $ start_station_name: chr [1:426887] "Western Ave & Leland Ave" "Clark St & Montrose Ave" "Broadway
## $ start_station_id : num [1:426887] 239 234 296 51 66 212 96 96 212 38 ...
## $ end_station_name : chr [1:426887] "Clark St & Leland Ave" "Southport Ave & Irving Park Rd" "Wilt
## $ end_station_id : num [1:426887] 326 318 117 24 212 96 212 212 96 100 ...
## $ start_lat : num [1:426887] 42 42 41.9 41.9 41.9 ...
## $ start_lng : num [1:426887] -87.7 -87.7 -87.6 -87.6 -87.6 ...
## $ end_lat : num [1:426887] 42 42 41.9 41.9 41.9 ...
## $ end_lng : num [1:426887] -87.7 -87.7 -87.7 -87.6 -87.6 ...
## $ member_casual : chr [1:426887] "member" "member" "member" "member" ...
## - attr(*, "spec")=
## .. cols(
## .. ride_id = col_character(),
## .. rideable_type = col_character(),
## .. started_at = col_datetime(format = ""),
```

```
## .. ended_at = col_datetime(format = ""),
## .. start_station_name = col_character(),
## .. start_station_id = col_double(),
## .. end_station_name = col_character(),
## .. end_station_id = col_double(),
## .. start_lat = col_double(),
## .. start_lng = col_double(),
## .. end_lat = col_double(),
## .. end_lng = col_double(),
## .. member_casual = col_character()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
str(q4_2019)
```

```
## spec_tbl_df [704,054 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ ride_id : num [1:704054] 25223640 25223641 25223642 25223643 25223644 ...
## $ started_at : POSIXct[1:704054], format: "2019-10-01 00:01:39" "2019-10-01 00:02:16" ...
## $ ended_at : POSIXct[1:704054], format: "2019-10-01 00:17:20" "2019-10-01 00:06:34" ...
## $ rideable_type : num [1:704054] 2215 6328 3003 3275 5294 ...
## $ tripduration : num [1:704054] 940 258 850 2350 1867 ...
## $ start_station_id : num [1:704054] 20 19 84 313 210 156 84 156 156 336 ...
## $ start_station_name: chr [1:704054] "Sheffield Ave & Kingsbury St" "Throop (Loomis) St & Taylor St"
## $ end_station_id : num [1:704054] 309 241 199 290 382 226 142 463 463 336 ...
## $ end_station_name : chr [1:704054] "Leavitt St & Armitage Ave" "Morgan St & Polk St" "Wabash Ave &
## $ member_casual : chr [1:704054] "Subscriber" "Subscriber" "Subscriber" "Subscriber" ...
## $ gender : chr [1:704054] "Male" "Male" "Female" "Male" ...
## $ birthyear : num [1:704054] 1987 1998 1991 1990 1987 ...
## - attr(*, "spec")=
## .. cols(
## .. trip_id = col_double(),
## .. start_time = col_datetime(format = ""),
## .. end_time = col_datetime(format = ""),
## .. bikeid = col_double(),
## .. tripduration = col_number(),
## .. from_station_id = col_double(),
## .. from_station_name = col_character(),
## .. to_station_id = col_double(),
## .. to_station_name = col_character(),
## .. usertype = col_character(),
## .. gender = col_character(),
## .. birthyear = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
str(q3_2019)
```

```
## spec_tbl_df [1,640,718 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ ride_id : num [1:1640718] 23479388 23479389 23479390 23479391 23479392 ...
## $ started_at : POSIXct[1:1640718], format: "2019-07-01 00:00:27" "2019-07-01 00:01:16" ...
## $ ended_at : POSIXct[1:1640718], format: "2019-07-01 00:20:41" "2019-07-01 00:18:44" ...
## $ rideable_type : num [1:1640718] 3591 5353 6180 5540 6014 ...
## $ tripduration : num [1:1640718] 1214 1048 1554 1503 1213 ...
```

```
## $ start_station_id : num [1:1640718] 117 381 313 313 168 300 168 313 43 43 ...
## $ start_station_name: chr [1:1640718] "Wilton Ave & Belmont Ave" "Western Ave & Monroe St" "Lakeview ...
## $ end_station_id : num [1:1640718] 497 203 144 144 62 232 62 144 195 195 ...
## $ end_station_name : chr [1:1640718] "Kimball Ave & Belmont Ave" "Western Ave & 21st St" "Larrabee ...
## $ member_casual : chr [1:1640718] "Subscriber" "Customer" "Customer" "Customer" ...
## $ gender : chr [1:1640718] "Male" NA NA NA ...
## $ birthyear : num [1:1640718] 1992 NA NA NA NA ...
## - attr(*, "spec")=
## .. cols(
## .. trip_id = col_double(),
## .. start_time = col_datetime(format = ""),
## .. end_time = col_datetime(format = ""),
## .. bikeid = col_double(),
## .. tripduration = col_number(),
## .. from_station_id = col_double(),
## .. from_station_name = col_character(),
## .. to_station_id = col_double(),
## .. to_station_name = col_character(),
## .. usertype = col_character(),
## .. gender = col_character(),
## .. birthyear = col_double()
## .. )
## - attr(*, "problems")=<externalptr>
```

```
str(q2_2019)
```

```
## spec_tbl_df [1,108,163 x 12] (S3: spec_tbl_df/tbl_df/tbl/data.frame)
## $ ride_id : num [1:1108163] 22178529 22178530 22178531 22178532 ...
## $ started_at : POSIXct[1:1108163], format: "2019-04-01 00:02:29" ...
## $ ended_at : POSIXct[1:1108163], format: "2019-04-01 00:09:49" ...
## $ rideable_type : num [1:1108163] 6251 6226 5649 4151 3270 ...
## $ 01 - Rental Details Duration In Seconds Uncapped: num [1:1108163] 446 1048 252 357 1007 ...
## $ start_station_id : num [1:1108163] 81 317 283 26 202 420 503 260 2 ...
## $ start_station_name : chr [1:1108163] "Daley Center Plaza" "Wood St & ...
## $ end_station_id : num [1:1108163] 56 59 174 133 129 426 500 499 2 ...
## $ end_station_name : chr [1:1108163] "Desplaines St & Kinzie St" "Wal ...
## $ member_casual : chr [1:1108163] "Subscriber" "Subscriber" "Subs ...
## $ Member Gender : chr [1:1108163] "Male" "Female" "Male" "Male" ...
## $ 05 - Member Details Member Birthday Year : num [1:1108163] 1975 1984 1990 1993 1992 ...
## - attr(*, "spec")=
## .. cols(
## .. '01 - Rental Details Rental ID' = col_double(),
## .. '01 - Rental Details Local Start Time' = col_datetime(format = ""),
## .. '01 - Rental Details Local End Time' = col_datetime(format = ""),
## .. '01 - Rental Details Bike ID' = col_double(),
## .. '01 - Rental Details Duration In Seconds Uncapped' = col_number(),
## .. '03 - Rental Start Station ID' = col_double(),
## .. '03 - Rental Start Station Name' = col_character(),
## .. '02 - Rental End Station ID' = col_double(),
## .. '02 - Rental End Station Name' = col_character(),
## .. 'User Type' = col_character(),
## .. 'Member Gender' = col_character(),
## .. '05 - Member Details Member Birthday Year' = col_double()
## .. )
```

```
## - attr(*, "problems")=<externalptr>
```

```
q4_2019 <- mutate(q4_2019, ride_id = as.character(ride_id)
                  ,rideable_type = as.character(rideable_type))
q3_2019 <- mutate(q3_2019, ride_id = as.character(ride_id)
                  ,rideable_type = as.character(rideable_type))
q2_2019 <- mutate(q2_2019, ride_id = as.character(ride_id)
                  ,rideable_type = as.character(rideable_type))
```

Converting ride_id and rideable_type to character so that they can stack correctly

```
all_trips <- bind_rows(q2_2019, q3_2019, q4_2019, q1_2020)
```

Stacking individual quarter's data frames into one big data frame

```
all_trips <- all_trips %>%
  select(-c(start_lat, start_lng, end_lat, end_lng, birthyear, gender, "01 - Rental Details Duration In
```

Removing lat, long, birthyear, and gender fields as this data was dropped starting from 2020

CLEAN UP AND ADD DATA FOR ANALYSIS

Let's begin by inspecting the table that has been created

```
colnames(all_trips) #List of column names
```

```
## [1] "ride_id"          "started_at"       "ended_at"
## [4] "rideable_type"    "start_station_id" "start_station_name"
## [7] "end_station_id"   "end_station_name" "member_casual"
```

```
nrow(all_trips) #How many rows are in data frame?
```

```
## [1] 3879822
```

```
dim(all_trips) #Dimensions of the data frame?
```

```
## [1] 3879822      9
```

```
head(all_trips) #See the first 6 rows of data frame.
```



```
## # A tibble: 6 x 9
##   ride_id started_at      ended_at      rideable_type start-1 start-2
##   <chr>    <dtm>        <dtm>        <chr>        <dbl> <chr>
## 1 22178529 2019-04-01 00:02:22 2019-04-01 00:09:48 6251      81 Daley ~
## 2 22178530 2019-04-01 00:03:02 2019-04-01 00:20:30 6226      317 Wood S~
## 3 22178531 2019-04-01 00:11:07 2019-04-01 00:15:19 5649      283 LaSall~
## 4 22178532 2019-04-01 00:13:01 2019-04-01 00:18:58 4151       26 McClur~
## 5 22178533 2019-04-01 00:19:26 2019-04-01 00:36:13 3270      202 Halste~
## 6 22178534 2019-04-01 00:19:39 2019-04-01 00:23:56 3123      420 Ellis ~
## # ... with 3 more variables: end_station_id <dbl>, end_station_name <chr>,
## #   member_casual <chr>, and abbreviated variable names 1: start_station_id,
## #   2: start_station_name
```

```
str(all_trips) #See list of columns and data types (numeric, character, etc)
```

```
## tibble [3,879,822 x 9] (S3: tbl_df/tbl/data.frame)
## $ ride_id      : chr [1:3879822] "22178529" "22178530" "22178531" "22178532" ...
## $ started_at   : POSIXct[1:3879822], format: "2019-04-01 00:02:22" "2019-04-01 00:03:02" ...
## $ ended_at     : POSIXct[1:3879822], format: "2019-04-01 00:09:48" "2019-04-01 00:20:30" ...
## $ rideable_type: chr [1:3879822] "6251" "6226" "5649" "4151" ...
## $ start_station_id : num [1:3879822] 81 317 283 26 202 420 503 260 211 211 ...
## $ start_station_name: chr [1:3879822] "Daley Center Plaza" "Wood St & Taylor St" "LaSalle St & Jack
## $ end_station_id   : num [1:3879822] 56 59 174 133 129 426 500 499 211 211 ...
## $ end_station_name  : chr [1:3879822] "Desplaines St & Kinzie St" "Wabash Ave & Roosevelt Rd" "Canal
## $ member_casual    : chr [1:3879822] "Subscriber" "Subscriber" "Subscriber" "Subscriber" ...
```

```
summary(all_trips) #Statistical summary of data; mainly for numeric data
```

```
##      ride_id      started_at
## Length:3879822  Min.   :2019-04-01 00:02:22.00
## Class :character 1st Qu.:2019-06-23 07:49:09.25
## Mode :character  Median :2019-08-14 17:43:38.00
##                      Mean  :2019-08-26 00:49:59.38
##                      3rd Qu.:2019-10-12 12:10:21.00
##                      Max.   :2020-03-31 23:51:34.00
##
##      ended_at      rideable_type      start_station_id
## Min.   :2019-04-01 00:09:48.00  Length:3879822  Min.   : 1.0
## 1st Qu.:2019-06-23 08:20:27.75  Class :character 1st Qu.: 77.0
## Median :2019-08-14 18:02:04.00  Mode  :character Median :174.0
## Mean   :2019-08-26 01:14:37.06                      Mean  :202.9
## 3rd Qu.:2019-10-12 12:36:16.75                      3rd Qu.:291.0
## Max.   :2020-05-19 20:10:34.00                      Max.   :675.0
##
##      start_station_name end_station_id end_station_name member_casual
## Length:3879822  Min.   : 1.0  Length:3879822  Length:3879822
## Class :character 1st Qu.: 77.0  Class :character  Class :character
## Mode :character  Median :174.0  Mode :character  Mode :character
##                      Mean   :203.8
##                      3rd Qu.:291.0
##                      Max.   :675.0
##                      NA's   :1
```

Now, we have discovered a few problems that we need to fix such as:

- In the “member_casual” column, there are two names for members (“member” and “Subscriber”) and two names for casual riders (“Customer” and “casual”). We will need to consolidate that from four to two labels.
- The data can only be aggregated at the ride-level, which is too granular. We will want to add some additional columns of data – such as day, month, year – that will provide additional opportunities to aggregate the data.
- We will want to add a calculated field for length of ride since the 2020_Q1 data did not have the “trip_duration” column. We will add “ride_length” to the entire dataframe for consistency.
- There are some rides where trip_duration shows up as negative, including several hundred rides where Divvy took bikes out of circulation for Quality Control reasons. We will want to delete these rides.

Way forward:

- In the “member_casual” column, we’ll replace “Subscriber” with “member” and “Customer” with “casual”.
- Before 2020, Divvy used different labels for these two types of riders ... we will want to make our dataframe consistent with their current nomenclature
- N.B.: “Level” is a special property of a column that is retained even if a subset does not contain any values from a specific level

```
table(all_trips$member_casual)
```

Begin by seeing how many observations fall under each usertype

```
##
##      casual    Customer      member Subscriber
##      48480      857474      378407    2595461
```

```
all_trips <- all_trips %>%
  mutate(member_casual = recode(member_casual
                                , "Subscriber" = "member"
                                , "Customer" = "casual"))
```

Reassign to the desired values (we will go with the current 2020 labels)

```
table(all_trips$member_casual)
```

Checking to make sure the proper number of observations were reassigned

```
##
##      casual  member
##    905954 2973868
```

Then add columns that list the date, month, day, and year of each ride. This will allow us to aggregate ride data for each month, day, or year ... before completing these operations we could only aggregate at the ride level

```
all_trips$date <- as.Date(all_trips$started_at) #The default format is yyyy-mm-dd
all_trips$month <- format(as.Date(all_trips$date), "%m")
all_trips$day <- format(as.Date(all_trips$date), "%d")
all_trips$year <- format(as.Date(all_trips$date), "%Y")
all_trips$day_of_week <- format(as.Date(all_trips$date), "%A")
```

Then we add a “ride_length” calculation to all_trips (in seconds)

```
all_trips$ride_length <- difftime(all_trips$ended_at, all_trips$started_at)
```

```
str(all_trips)
```

Inspecting the structure of the columns

```
## tibble [3,879,822 x 15] (S3: tbl_df/tbl/data.frame)
##  $ ride_id          : chr [1:3879822] "22178529" "22178530" "22178531" "22178532" ...
##  $ started_at       : POSIXct[1:3879822], format: "2019-04-01 00:02:22" "2019-04-01 00:03:02" ...
##  $ ended_at         : POSIXct[1:3879822], format: "2019-04-01 00:09:48" "2019-04-01 00:20:30" ...
##  $ rideable_type     : chr [1:3879822] "6251" "6226" "5649" "4151" ...
##  $ start_station_id : num [1:3879822] 81 317 283 26 202 420 503 260 211 211 ...
##  $ start_station_name: chr [1:3879822] "Daley Center Plaza" "Wood St & Taylor St" "LaSalle St & Jack
##  $ end_station_id    : num [1:3879822] 56 59 174 133 129 426 500 499 211 211 ...
##  $ end_station_name  : chr [1:3879822] "Desplaines St & Kinzie St" "Wabash Ave & Roosevelt Rd" "Canal
##  $ member_casual     : chr [1:3879822] "member" "member" "member" "member" ...
##  $ date              : Date[1:3879822], format: "2019-04-01" "2019-04-01" ...
##  $ month             : chr [1:3879822] "04" "04" "04" "04" ...
##  $ day               : chr [1:3879822] "01" "01" "01" "01" ...
##  $ year              : chr [1:3879822] "2019" "2019" "2019" "2019" ...
##  $ day_of_week       : chr [1:3879822] "Monday" "Monday" "Monday" "Monday" ...
##  $ ride_length       : 'difftime' num [1:3879822] 446 1048 252 357 ...
##  ..- attr(*, "units")= chr "secs"
```

```
is.factor(all_trips$ride_length)
```

Converting “ride_length” from Factor to numeric so we can run calculations on the data

```
## [1] FALSE
```

```
all_trips$ride_length <- as.numeric(as.character(all_trips$ride_length))
is.numeric(all_trips$ride_length)
```

```
## [1] TRUE
```

Removing “bad” data The dataframe includes a few hundred entries when bikes were taken out of docks and checked for quality by Divvy or ride_length was negative. We will create a new version of the dataframe (v2) since data is being removed.

```
all_trips_v2 <- all_trips[!(all_trips$start_station_name == "HQ QR" | all_trips$ride_length<0),]
```

CONDUCT DESCRIPTIVE ANALYSIS

Descriptive analysis on ride_length (all figures in seconds). Using the summary(), we can get the mean/straight average (which is; total ride length / rides), median (midpoint number in the ascending array of ride lengths), max (longest ride), min (shortest ride).

```
summary(all_trips_v2$ride_length)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.     Max.
##         1      412      712    1479    1289 9387024
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = mean)
```

Comparing members and casual users

```
##   all_trips_v2$member_casual all_trips_v2$ride_length
## 1                        casual      3552.7502
## 2                        member       850.0662
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = median)
```

```
##   all_trips_v2$member_casual all_trips_v2$ride_length
## 1                        casual           1546
## 2                        member            589
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = max)
```

```
##   all_trips_v2$member_casual all_trips_v2$ride_length
## 1                        casual      9387024
## 2                        member     9056634
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual, FUN = min)
```

```
##   all_trips_v2$member_casual all_trips_v2$ride_length
## 1                        casual                2
## 2                        member                1
```

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual + all_trips_v2$day_of_week, FUN = mean)
```

Viewing the average ride time by each day for members vs casual users

```
##      all_trips_v2$member_casual all_trips_v2$day_of_week all_trips_v2$ride_length
## 1          casual          Friday          3773.8351
## 2          member          Friday           824.5305
## 3          casual          Monday          3372.2869
## 4          member          Monday           842.5726
## 5          casual          Saturday         3331.9138
## 6          member          Saturday           968.9337
## 7          casual          Sunday          3581.4054
## 8          member          Sunday           919.9746
## 9          casual          Thursday         3682.9847
## 10         member          Thursday           823.9278
## 11         casual          Tuesday          3596.3599
## 12         member          Tuesday           826.1427
## 13         casual          Wednesday         3718.6619
## 14         member          Wednesday           823.9996
```

```
all_trips_v2$day_of_week <- ordered(all_trips_v2$day_of_week, levels=c("Sunday", "Monday", "Tuesday", "
```

We may Notice that the days of the week are out of order. Let's fix that.

```
aggregate(all_trips_v2$ride_length ~ all_trips_v2$member_casual + all_trips_v2$day_of_week, FUN = mean)
```

Running the average ride time by each day for members vs casual users again:

```
##      all_trips_v2$member_casual all_trips_v2$day_of_week all_trips_v2$ride_length
## 1          casual          Sunday          3581.4054
## 2          member          Sunday           919.9746
## 3          casual          Monday          3372.2869
## 4          member          Monday           842.5726
## 5          casual          Tuesday          3596.3599
## 6          member          Tuesday           826.1427
## 7          casual          Wednesday         3718.6619
## 8          member          Wednesday           823.9996
## 9          casual          Thursday         3682.9847
## 10         member          Thursday           823.9278
## 11         casual          Friday          3773.8351
## 12         member          Friday           824.5305
## 13         casual          Saturday         3331.9138
## 14         member          Saturday           968.9337
```

```
all_trips_v2 %>%
  mutate(weekday = wday(started_at, label = TRUE)) %>% #creates weekday field using wday()
  group_by(member_casual, weekday) %>% #groups by usertype and weekday
  summarise(number_of_rides = n() #calculates the number of rides and average
            ,average_duration = mean(ride_length)) %>% # calculates the average duration
  arrange(member_casual, weekday)
```

analyzing ridership data by type and weekday


'summarise()' has grouped output by 'member_casual'. You can override using the
'.groups' argument.

```
## # A tibble: 14 x 4
## # Groups:   member_casual [2]
##   member_casual weekday number_of_rides average_duration
##   <chr>          <ord>          <int>          <dbl>
## 1 casual        Sun            181293         3581.
## 2 casual        Mon            103296         3372.
## 3 casual        Tue             90510         3596.
## 4 casual        Wed             92457         3719.
## 5 casual        Thu            102679         3683.
## 6 casual        Fri            122404         3774.
## 7 casual        Sat            209543         3332.
## 8 member        Sun             267965           920.
## 9 member        Mon             472196           843.
## 10 member       Tue             508445           826.
## 11 member       Wed             500329           824.
## 12 member       Thu             484177           824.
## 13 member       Fri             452790           825.
## 14 member       Sat             287958           969.
```

```
all_trips_v2 %>%
  mutate(weekday = wday(started_at, label = TRUE)) %>%
  group_by(member_casual, weekday) %>%
  summarise(number_of_rides = n()
            ,average_duration = mean(ride_length)) %>%
  arrange(member_casual, weekday) %>%
  ggplot(aes(x = weekday, y = number_of_rides, fill = member_casual)) +
  geom_col(position = "dodge")
```

visualizing the number of rides by rider type

'summarise()' has grouped output by 'member_casual'. You can override using the
'.groups' argument.



```
all_trips_v2 %>%
  mutate(weekday = wday(started_at, label = TRUE)) %>%
  group_by(member_casual, weekday) %>%
  summarise(number_of_rides = n(),
            ,average_duration = mean(ride_length)) %>%
  arrange(member_casual, weekday) %>%
  ggplot(aes(x = weekday, y = average_duration, fill = member_casual)) +
  geom_col(position = "dodge")
```

Creating visualization for average duration

'summarise()' has grouped output by 'member_casual'. You can override using the
'.groups' argument.

